



Data Warehouses

Lecture 13

Lecture Goals

Goal:

Physical aspects of OLAP


materialization

indexing,

partitioning,

columnar storage

Physical aspects



What is Physical Design?

Physical design is the phase following logical design that identifies actual database tables and index structures used to implement the data warehouse.

This phase translates logical models into implementable database structures, making critical decisions that influence overall performance and maintenance.

Physical design decisions are primarily driven by performance optimization and operational efficiency requirements.

The physical design of the data warehouses is crucial to ensure adequate query response time

underlying data storage and processing solution

There are typically three common techniques for improving performance in data warehouse systems:

Materialized view

is a view that is physically stored in a database, which enhances query performance by precalculating costly operations such as joins and aggregations.

Indexing

alternative indexing mechanisms are used in data warehouses, typically bitmap and join indexes.

Partitioning or fragmentation

divides the contents of a relation into several files, typically based on a range of values of an attribute

Materialized views enhance query performance

by precalculating costly operations such as joins and aggregations achieved at the expense of storage space.

Typical problems related to the use of materialized views:

Selection

select an appropriate set of views that minimizes the total query response time and the cost of maintaining the selected views

Calculation

Efficient creation of materialized views

Update

since all modifications to the underlying base tables must be propagated into the view (view maintenance)

Query rewrite

Trying to use the materialized views as much as possible, even if they only partially fulfill the query conditions

Since data warehouse queries typically access a large number of tuples

alternative indexing mechanisms are needed

Two common types of indexes for data warehouses are

Bitmap indexes – useful for columns with a low number of distinct values (i.e., low cardinality attributes)

several compression techniques can overcome low cardinality limitation

Bitslice Indices – use of bitmap indices by using bitmaps to index value ranges

Join indexes – materialize a relational join between two tables by keeping pairs of row identifiers that participate in the join.

In data warehouses, join indexes relate the values of dimensions to rows in the fact table.

Star index – join index between a fact table and each of its dimension tables

Join indexes can be combined with bitmap indexes - bitmap join indexes

Store pre-computed join results between fact and dimension tables - pre-materialized join results

Other options

aggregate index, summary index, hybrid b-tree/bitmap

Partitioning or fragmentation is a mechanism frequently used in relational databases to reduce the execution time of queries.

It consists in dividing the contents of a relation into several files that can be more efficiently processed in this way.

There are two ways of partitioning a relation:

Vertical partitioning – splits the attributes of a table into groups that can be independently stored.

For example, a table can be partitioned such that the most often used attributes are stored in one partition, while other less often used attributes are kept in another partition.

Horizontal partitioning – divides the records of a table into groups according to a particular criterion.

A common horizontal partitioning scheme in data warehouses is based on time, where each partition contains data about a particular time period, for instance, a year or a range of months.

Columnar storage

column store database is a type of database that stores data using a column-oriented model („stores all column data together”) – vertical partitioning

- Significantly increases performance for OLAP queries

- Allows for better compression

has become the backbone of modern data warehousing due to its exceptional ability to handle analytical workloads efficiently.

- Data warehouses primarily serve analytical workloads that scan large numbers of rows while typically touching only a small subset of columns


- Unlike traditional row-based storage systems, columnar storage organizes data by columns rather than rows, creating significant advantages for the types of queries commonly performed in data warehouse environments

 - In a columnar data warehouse, each column of a table is stored separately as a contiguous block of data on disk

- Benefits are substantial

 - Using columnar storage, each data block can hold column field values for as many as three times the number of records compared to row-based storage, reducing I/O operations by approximately two-thirds


Materialization



Aggregates

In multidimensional data model together with measure values usually we store summarizing information

aggregates



Location: Vancouver					
Time (quarters)	Items				
	TV	Computer	Phone	Security	Sum
Q1	605	825	14	400	1844
Q2	680	952	31	512	2175
Q3	812	1023	30	501	2366
Q4	927	1038	38	580	2583
Sum	3024	3838	113	1993	8968

Materialized views

Materialized views that hold aggregates

are one of the most important means of improving query performance in a multidimensional database management system

Precomputed aggregate view is the materialized result of a query


that aggregates (e.g., summarizes) parts of or all the data in the database

The word “materialized” means that the result is stored physically

the system can answer the query by just reading the (already computed and stored) result


instead of computing it from the detailed base data

Cube Aggregates Lattice



Aggregates

In multidimensional data model together with measure values we store summarizing information aggregates



Consider the following 2 cases for n-dimensional cube

Case 1: Dimensions have no hierarchies

Then the total number of cuboids computed for a n-dimensional cube

$$2^n$$

Case 2: Dimensions have hierarchies

Then the total number of cuboids computed for a n-dimensional cube

$$\prod_{i=1}^n (L_i + 1)$$

where L_i is the number of levels associated with dimension i ; n is the number of dimensions

Materializing cubes

Example

Cube with 3 dimensions – item,
city, year

2^3 cuboids


Need compute the following
groupings (group by)

(item, city, year),

(item, city), (item, year), (city, year),

(item), (city), (year)

()



Materializing cubes

Example


Cube with 2 dimensions

One dimension with 4 levels

One dimension with 3 levels


In total:

$$5 \times 4 = 20$$



Materializing cubes

Why cuboids are important?



Processing

Materialized views

The use of materialized views can speed up query processing

to be several hundreds or thousands times faster than when using the base data.

Often many different queries can even be answered efficiently by using the same materialized view.

If we, for example, have an aggregate that holds the summarized sales for cities, we can also use that aggregate to calculate the total sales for states and the single total of all sales.

Transparent to the end user

should not be burdened with the task of deciding when to use a precomputed aggregate

OLAP system must take care of this transparently to the user

Materializing cubes

View materialization in data warehouses presents a fundamental trade-off between three competing factors:

*query processing cost,
view maintenance cost,
and storage space requirements*

The goal is to select an appropriate set of views

that minimizes the sum of query response time and maintenance costs while operating within resource constraints such as storage space and maintenance time window

Materializing cubes

The problem is then to decide which precomputed aggregates to create.

One extreme is to not have any.

But if there are no precomputed aggregates, query responses will be slow.

The other extreme is to precompute each possible aggregate.

Unfortunately, they will occupy huge amounts of disk space and will take unreasonably long to compute in the first place.

Further, they need to be updated when updates to the base data occur.

So this extreme is also not practical.

A good compromise in-between the two extremes must therefore be found.

Unfortunately, the problem of picking the optimal set of aggregates is an NP hard problem.

Materializing cubes

The problem lies in selection of cuboids to be materialized

high number of materialized cuboids


huge size of data warehouse.

small number of materialized cuboids

slow query processing

use information to balance the two

like size and storage constraints, dependency, update and maintenance costs, access frequency and patterns, cost (mixture of the former)



Materializing cubes

Materialization of data cube

Materialize every cuboid (full materialization)

full cube

all the cells of all of the cuboids for a given data cube

Materialize none (no materialization)

base cube

minimal set of cells cube

Materialize some (partial materialization)

partial cube

subset of full cube

Partial

Partial cube materialization

Selected parts of a data cube (selected cuboids) are materialized

Hybrid approach

Balance between the response time and required storage space and overhead

Instead of computing the full cube, we can compute only a subset of the data cube's cuboids, or subcubes consisting of subsets of cells from the various cuboids.

Requires

to identify a subset of cuboids that will be materialized

Many cells in a cuboid may actually be of little or no interest to the data analyst

efficient update during each data load and refresh

We can use some metadata

Usage profiles, Sparsity information, Etc.

Partial

Performance vs. Space Trade –off

Maximum performance boost


implies using lots of disk space for storing every pre calculation.

Minimum performance boost

implies no disk space with zero pre - calculation.

Using meta-data/statistics

to determine best level of pre-aggregation from which all other aggregates can be computed.



Materializing cubes

A typical problem of materialized views is

updating

since all modifications to the underlying base tables must be propagated into the view (view maintenance)

selection of materialized views

select an appropriate set of views that minimizes the total query response time and the cost of maintaining the selected views

query rewriting

tries to use the materialized views as much as possible, even if they only partially fulfill the query conditions

Materialized View Selection

Several OLAP products have adopted heuristic approaches for cuboid and subcube selection.

A popular approach is to materialize the set of cuboids on which other frequently referenced cuboids are based.

e.g. Greedy Lattice-Based Selection (The HRU Algorithm)

Constraint-Based / Resource-Aware Materialization

*Alternatively, we can compute an **iceberg cube***

data cube that stores only those cube cells whose aggregate value (e.g., count) is above some minimum support threshold.

*Another common strategy is to materialize a **shell cube**.*

This involves precomputing the cuboids for only a small number of dimensions (such as 3 to 5) of a data cube.

*The dataset's dimensions are partitioned into disjoint groups and each group, called a **shell fragment**, is processed independently - **full local data cube** is computed for each fragment*

retains inverted indices mapping attribute values to tuple IDs

Use usage information – Workload-Directed / Query-Directed Selection

Materializing cubes

Effective view materialization relies on comprehensive cost models

that evaluate the total expenses associated with each materialization decision

These models typically incorporate:

Query Processing Costs:

The computational expense of executing queries against the selected set of materialized views, considering factors such as join operations, aggregations, and data scanning requirements.

Maintenance Costs:

The overhead associated with keeping materialized views synchronized with their underlying base tables, including the time and resources required for incremental or complete refreshes.

Storage Costs:

The disk space requirements for storing materialized views, which directly impacts both hardware costs and system performance.

Warehouse Maintenance

Warehouse data $\approx$ materialized view

Initial loading

View maintenance

Derived Warehouse Data

indexes

aggregates

materialized views

View maintenance



Choosing view

Query rewriting

Materialized views

Efficient Processing of OLAP Queries

The purpose of materializing cuboids and constructing OLAP index structures is to speed up query processing in data cubes.

Given materialized views, query processing should proceed as follows:

1. Determine which operations should be performed on the available cuboids

This involves transforming any selection, projection, roll-up (group-by), and drill-down operations specified in the query into corresponding SQL and/or OLAP operations.

For example, slicing and dicing a data cube may correspond to selection and/or projection operations on a materialized cuboid.

2. Determine to which materialized cuboid(s) the relevant operations should be applied

Choose view to answer a query

Identify all the materialized cuboids that may potentially be used to answer a query.

Pruning the above set using knowledge of dominance relationship among cuboids.

Estimating the costs of using the remaining materialized cuboids and selecting the cuboid with the least cost.

Example

Cube

Fact Table - sales [time, item, location]:

sum(sales_in_dollars)

Dimension's hierarchies used are :

day<month<quarter<year for time

item_name<brand<type for item

street<city<province_or_state<country for location

Cuboids are :

cuboid 1: {item, city, year}

cuboid 2: {brand, country, year}

cuboid 3: {brand, state, year}

cuboid 4: {item, state} where year = 2000

Query

{brand, state} with year = 2000?

Analysis

cuboid 1

costs the most since item and city are at lower level than brand and state

cuboid 2

can not be used since state < country

cuboid 3

is better than cuboid 4 if there are few year values associated with items and several items for each brand

cuboid 4

is better than cuboid 3 if efficient indices are available for cuboid 4.

Example

cuboid 1:

{item, city, year}

cuboid 2:

{brand, country, year}

cuboid 3:

{brand, state, year}

cuboid 4:

{item, state} where year = 2000

Query

{brand, state} with year = 2000?

What is the optimal set of views to materialize?

How many impressions do we need to materialize to get reasonable performance?

(part, supplier)
 SELECT part, supplier, sum(sales)
 FROM Sales
 GROUP BY part, supplier

(supplier)
 SELECT supplier, sum(sales)
 FROM Sales
 GROUP BY supplier

(part)
 SELECT part, sum(sales)
 FROM Sales
 GROUP BY part

(part, customer)
 SELECT part, customer, sum(sales)
 FROM Sales
 GROUP BY part, customer

Query	Answer	Cost
• (part,supplier,customer)	6M	6M
• (part,customer)	6M	6M
• (part,supplier)	0.8M	6M
• (supplier,customer)	6M	6M
• (part)	0.2M	6M
• (supplier)	0.01M	6M
• (customer)	0.1M	6M
• ()	1	6M

(part, supplier)
 SELECT part, supplier, sum(sales)
 FROM Sales
 GROUP BY part, supplier

(supplier)
 SELECT supplier, sum(sales)
 FROM Sales
 GROUP BY supplier

(part)
 SELECT part, sum(sales)
 FROM Sales
 GROUP BY part

(part, customer)
 SELECT part, customer, sum(sales)
 FROM Sales
 GROUP BY part, customer

Query	Answer	Cost
• <u>(part,supplier,customer)</u>	6M	<u>6M</u>
• (part,customer)	6M	6M
• <u>(part,supplier)</u>	0.8M	<u>0.8M</u>
• (supplier,customer)	6M	6M
• (part)	0.2M	<u>0.8M</u>
• (supplier)	0.01M	<u>0.8M</u>
• <u>(customer)</u>	0.1M	<u>0.1M</u>
• ()	1	<u>0.1M</u>



Indexing in DW

Indexing is another important means of improving query performance

Fact tables are typically very large and are used in the majority of queries.

Therefore, a fact table often has many indices such that most queries can benefit from one or more of them

Typically, a separate index is built on each dimension key

If the DBMS supports index intersection, the indices can then be used in combination when answering a query

It can also be a good idea, however, to create indices on combinations of dimension keys for those combinations that are often used together by queries

In general, the key referencing the Date dimension should be put first in combinations since most queries refer to dates

Common approach:

- Clustered / Sort Keys (On Date/Time)

- Zone Maps / Data Skipping (On Sort Keys and Measures)

- Bitmap Indexes (On Foreign Keys)

Indexing is another important means of improving query performance

Dimension tables should also be indexed.

Depending on typical usage, it may or may not make sense to build indices on all individual columns.

As for fact tables, indices can be built on combinations of attributes.

Common approach:

- Standard B-Tree Indexes (On Primary Keys)

- Bitmap Indexes (On Descriptive Columns)

B+Tree

B-tree stores data such that each node contains keys in ascending order.

Each of these keys has two references to another two child nodes.


The left-side child node keys are less than the current keys,

and the right-side child node keys are more than the current keys.


If a single node has “n” number of keys then it can have maximum “n+1” child nodes.

B+tree stores data on the leaf nodes

This means that all non-leaf node values are duplicated in leaf nodes again.



Clustered




Sr. No.	Operation	Time Complexity
1.	Search	$O(\log n)$
2.	Insert	$O(\log n)$
3.	Delete	$O(\log n)$
4.	Traverse	$O(n)$

Note: "n" is the total number of elements in the B-tree


Join index

A join index:

an index on one table that involves a column value from different table through a commonly encountered join.



np. indeks na atrybucie Sklepy.sklep_id



Bitmaps

The principle of the bitmap index is that the position bitmap for the value v

has a 1 on position p

if and only if the p 'th row in the table takes the value v

has a 0

otherwise

A bitmap index makes it very fast to locate rows with a certain value

Just identify 1s

It is also possible to combine bitmap indices

Multiple indices is just computing the logical AND or OR or XOR

It is also possible to use other operators on bitmap indices

Negation

kwota	miasto
950,00	Warszawa
800,00	Kraków
755,00	Poznań
320,00	Kraków
110,00	Kraków
230,00	Warszawa
170,50	Poznań
190,00	Warszawa

Warszawa	Kraków	Poznań
1	0	0
0	1	0
0	0	1
0	1	0
0	1	0
1	0	0
0	0	1
1	0	0

Using Bitmaps

Query:

Get people with age = 20 and name = "fred"

Result:

Bit Map for age = 20:

1 1 0 1 1 0 0 0

Bit Map for name = "fred":

0 1 0 0 0 0 0 0

Answer is intersection:

_ 1 _ _ _ _ _

rld	name	age
r4	joe	20
r18	fred	20
r19	sally	21
r34	nancy	20
r35	tom	20
r36	pat	25
r5	dave	21
r41	jeff	26

Bitmap Index

Index on a particular column

Each value in the column has a bit vector: bit-op is fast

The length of the bit vector: # of records in the base table

The i -th bit is set if the i -th row of the base table has the value for the indexed column

not suitable for high cardinality domains

Base table

Cust	Region	Type
C1	Asia	Retail
C2	Europe	Dealer
C3	Asia	Dealer
C4	America	Retail
C5	Europe	Dealer

Index on Region

RecID	Asia	Europe	America
1	1	0	0
2	0	1	0
3	1	0	0
4	0	0	1
5	0	1	0

Index on Type

RecID	Retail	Dealer
1	1	0
2	0	1
3	0	1
4	1	0
5	0	1

Bitmap index

a

ProductKey	ProductName	QuantityPerUnit	UnitPrice	Discontinued	CategoryKey
p1	prod1	25	60	No	c1
p2	prod2	45	60	Yes	c1
p3	prod3	50	75	No	c2
p4	prod4	50	100	Yes	c2
p5	prod5	50	120	No	c3
p6	prod6	70	110	Yes	c4

b

	25	45	50	70
p1	1	0	0	0
p2	0	1	0	0
p3	0	0	1	0
p4	0	0	1	0
p5	0	0	1	0
p6	0	0	0	1

c

	60	75	100	110	120
p1	1	0	0	0	0
p2	1	0	0	0	0
p3	0	1	0	0	0
p4	0	0	1	0	0
p5	0	0	0	0	1
p6	0	0	0	1	0

Bitmap join index

In addition to a bitmap index on a single table, you can create a bitmap join index

a bitmap index for the join of two or more tables.

Bitmap join index

the bitmap for the table to be indexed is built for values coming from the joined tables.

In a data warehousing environment, the join condition is an equi-inner join between the primary key column or columns of the dimension tables and the foreign key column or columns in the fact table.

can improve the performance by an order of magnitude.

By storing the result of a join, the join can be avoided completely for SQL statements using a bitmap join index.

Bitmap join index

RowID Product	Product Key	Product Name	...	Discontinued	...
1	p1	prod1	...	No	...
2	p2	prod2	...	Yes	...
3	p3	prod3	...	No	...
4	p4	prod4	...	Yes	...
5	p5	prod5	...	No	...
6	p6	prod6	...	Yes	...

RowID Sales	Product Key	Customer Key	Time Key	Sales Amount
1	p1	c1	t1	100
2	p1	c2	t1	100
3	p2	c2	t2	100
4	p2	c2	t3	100
5	p3	c3	t3	100
6	p4	c3	t4	100
7	p5	c4	t5	100

join
index

RowID Sales	RowID Product
1	1
2	1
3	2
4	2
5	3
6	4
7	5

join bitmap
index

Yes	No
0	1
0	1
1	0
1	0
0	1
1	0
0	1

Bitmap join index

RowID Product	Product Key	Product Name	...	Discontinued	...
1	p1	prod1	...	No	...
2	p2	prod2	...	Yes	...
3	p3	prod3	...	No	...
4	p4	prod4	...	Yes	...
5	p5	prod5	...	No	...
6	p6	prod6	...	Yes	...

Customer Key	Customer Name	Address	Postal Code	...
c1	cust1	35 Main St.	7373	...
c2	cust2	Av. Roosevelt 50	1050	...
c3	cust3	Av. Louise 233	1080	...
c4	cust4	Rue Gabrielle	1180	...

Star queries

Require join of fact and dimension tables

Example

Total sales of discontinued products by client and product:

```
SELECT    C.CustomerName, P.ProductName, SUM(S.SalesAmount)
FROM      Sales S, Customer C, Product P
WHERE     S.CustomerKey = C.CustomerKey AND
          S.ProductKey = P.ProductKey AND P.Discontinued = 'Yes'
GROUP BY C.CustomerName, P.ProductName
```

How is this query handled by the system?

RowID Sales	Product Key	Customer Key	Time Key	Sales Amount
1	p1	c1	t1	100
2	p1	c2	t1	100
3	p2	c2	t2	100
4	p2	c2	t3	100
5	p3	c3	t3	100
6	p4	c3	t4	100
7	p5	c4	t5	100

B⁺-tree

Product Key	Product Name	...	Discontinued	...
p1	prod1	...	No	...
p2	prod2	...	Yes	...
p3	prod3	...	No	...
p4	prod4	...	Yes	...
p5	prod5	...	No	...
p6	prod6	...	Yes	...

Yes	No
0	1
1	0
0	1
1	0
0	1
1	0

Product Key	Customer Key	Time Key	Sales Amount
p1	c1	t1	100
p1	c2	t1	100
p2	c2	t2	100
p2	c2	t3	100
p3	c3	t3	100
p4	c3	t4	100
p5	c4	t5	100

p1	p2	p3	p4	p5	p6
1	0	0	0	0	0
1	0	0	0	0	0
0	1	0	0	0	0
0	1	0	0	0	0
0	0	1	0	0	0
0	0	0	1	0	0
0	0	0	0	1	0

c1	c2	c3	c4
1	0	0	0
0	1	0	0
0	1	0	0
0	1	0	0
0	0	1	0
0	0	1	0
0	0	0	1

Customer Key	Customer Name	Address	Postal Code	...
c1	cust1	35 Main St.	7373	...
c2	cust2	Av. Roosevelt 50	1050	...
c3	cust3	Av. Louise 233	1080	...
c4	cust4	Rue Gabrielle	1180	...

bitmap indexes

B⁺-tree

Product Key	Product Name	...	Discontinued	...
p1	prod1	...	No	...
p2	prod2	...	Yes	...
p3	prod3	...	No	...
p4	prod4	...	Yes	...
p5	prod5	...	No	...
p6	prod6	...	Yes	...

Product Key	Customer Key	Time Key	Sales Amount
p1	c1	t1	100
p1	c2	t1	100
p2	c2	t2	100
p2	c2	t3	100
p3	c3	t3	100
p4	c3	t4	100
p5	c4	t5	100

Yes	No
0	1
0	1
1	0
1	0
0	1
1	0
0	1

bitmap join indexes.


Customer Key	Customer Name	Address	Postal Code	...
c1	cust1	35 Main St.	7373	...
c2	cust2	Av. Roosevelt 50	1050	...
c3	cust3	Av. Louise 233	1080	...
c4	cust4	Rue Gabrielle	1180	...

Bitmap join index

id_produktu	nazwa_produktu
1	Jumanji
2	Hero
3	The Client

suma_sprzedazy	id_sklepu	id_produktu
5820,00	1	1
17200,00	1	2
350,00	2	2

id_sklepu	miestowosc
1	New Orleans
2	Chicago





Columnar storage

Columnar storage

We store the database in storage.

Data is written to storage in contiguous chunks.

In row oriented databases

these chunks contain a row consisting of all of its column values as a tuple.

In column oriented databases

these chunks contain only the values in each row that belongs to that column.

Given a certain amount of space on an SSD or HDD, a chunk of rows/columns data is read at once.

This difference has a significant impact on performance.

If the data you need to access is stored mostly in a small number of columns and it is not necessary to query each field in the rows

you may be better off with a columnar-store

If you need many columns in each row to determine which rows are relevant

a row-store may be a better fit.

Columnar storage

Example

Retrieving all columns from row 1 through row 5 requires a relatively low effort as the data is contiguous and ordered

Row 1-1	Row 1-2	Row 1-3	Row 1-4	Row 1-5	Row 2-1	Row 2-2	Row 2-3	Row 2-4
Row 2-5	Row 3-1	Row 3-2	Row 3-3	Row 3-4	Row 3-5	Row 4-1	Row 4-2	Row 4-3
Row 4-4	Row 4-5	Row 5-1	Row 5-2	Row 5-3	Row 5-4	Row 5-5	Row 6-1	Row 6-2
Row 6-3	Row 6-4	Row 6-5	Row 7-1	Row 7-2	Row 7-3	Row 7-4	Row 7-5	Row 8-1
Row 8-2	Row 8-3	Row 8-4	Row 8-5	Row 9-1	Row 9-2	Row 9-3	Row 9-4	Row 9-5

Consider a typical OLAP query against the example table where a single column is aggregated across a large portion of the table that happens to include all of the rows shown here.

Row 1-1	Row 1-2	Row 1-3	Row 1-4	Row 1-5	Row 2-1	Row 2-2	Row 2-3	Row 2-4
Row 2-5	Row 3-1	Row 3-2	Row 3-3	Row 3-4	Row 3-5	Row 4-1	Row 4-2	Row 4-3
Row 4-4	Row 4-5	Row 5-1	Row 5-2	Row 5-3	Row 5-4	Row 5-5	Row 6-1	Row 6-2
Row 6-3	Row 6-4	Row 6-5	Row 7-1	Row 7-2	Row 7-3	Row 7-4	Row 7-5	Row 8-1
Row 8-2	Row 8-3	Row 8-4	Row 8-5	Row 9-1	Row 9-2	Row 9-3	Row 9-4	Row 9-5

Columnar storage

Row-stores

Writing one row at a time is easy for a row-store because it appends the whole row to a chunk of space in storage.

the row oriented database is partitioned horizontally

Are also good for point lookups and index range scans.


Indexing (creating a key from columns) based on your access patterns can optimize queries and prevent full table scans in a large row-store

Note that creating many indices will create many copies of data

If only one field of the record is desired

using a row-store becomes expensive since all the fields in each record will be read.

Even data that isn't needed for the query response will be read, assuming it isn't indexed properly.



Columnar storage

Consider a different storage for a typical OLAP query

a single column is aggregated across a large portion of the table that happens to include all of the rows shown here.

Row 1-1	Row 2-1	Row 3-1	Row 4-1	Row 5-1	Row 6-1	Row 7-1	Row 8-1	Row 9-1
Row 1-2	Row 2-2	Row 3-2	Row 4-2	Row 5-2	Row 6-2	Row 7-2	Row 8-2	Row 9-2
Row 1-3	Row 2-3	Row 3-3	Row 4-3	Row 5-3	Row 6-3	Row 7-3	Row 8-3	Row 9-3
Row 1-4	Row 2-4	Row 3-4	Row 4-4	Row 5-4	Row 6-4	Row 7-4	Row 8-4	Row 9-4
Row 1-5	Row 2-5	Row 3-5	Row 4-5	Row 5-5	Row 6-5	Row 7-5	Row 8-5	Row 9-5

Columnar storage

Column-stores

Analytical applications often need aggregate data, where only a subset of a table's attributes are needed.

Individual values are stored contiguously in storage by column


Column oriented databases are partitioned vertically

The advantage of a columnar-store is that *partial reads are much more efficient*

because a lower volume of data is loaded due to reading only the relevant data instead of the whole record.

compression is achieved more efficiently

than in row-stores because columns have uniform types (ex: all strings or integers).



Columnar storage


SQL Server Page: Row Data

Row1Col1	Row1Col2	Row1Col3	Row1Col4	Row1Col5
Row1Col6	Row2Col1	Row2Col2	Row2Col3	Row2Col4
Row2Col5	Row2Col6	Row3Col1	Row3Col2	Row3Col3
Row3Col4	Row3Col5	Row3Col6	Row4Col1	Row4Col2
Row4Col3	Row4Col4	Row4Col5	Row4Col6	Row5Col1
Row5Col2	Row5Col3	Row5Col4	Row5Col5	Row5Col6
Row6Col1	Row6Col2	Row6Col3	Row6Col4	Row6Col5
Row6Col6	Row7Col1	Row7Col2	Row7Col3	Row7Col4

SQL Server Page: Column Data

Row1Col1	Row2Col1	Row3Col1	Row4Col1	Row5Col1
Row6Col1	Row7Col1	Row8Col1	Row9Col1	Row10Col1
Row11Col1	Row12Col1	Row13Col1	Row14Col1	Row15Col1
Row16Col1	Row17Col1	Row18Col1	Row19Col1	Row20Col1
Row21Col1	Row22Col1	Row23Col1	Row24Col1	Row25Col1
Row26Col1	Row27Col1	Row28Col1	Row29Col1	Row30Col1
Row31Col1	Row32Col1	Row33Col1	Row34Col1	Row35Col1
Row36Col1	Row37Col1	Row38Col1	Row39Col1	Row40Col1

Hybrid



Columnar storage

Example

Schema stores four columns in a table with 1 million rows.

Timestamp

Temperature

Humidity

Wind speed

Randomly generated data

Page-based Storage:

Organizes data in 4096-byte pages to optimize read/write efficiency.

Illustrative Results

- Fewer pages are read from disk with columnar storage

```
Filter Selectivity: 0.003416
Row Storage Query Result: 7358
Row Storage Query Time: 0.112099 seconds
Columnar Storage Query Result: 7358
Columnar Storage Query Time: 0.00361763 seconds
```

```
Pages Read:
Total Row Storage Pages Read: 3907
Filter Pages Read: 489
Aggregate Pages Read: 473
Total Columnar Storage Pages Read: 962
```

Column-oriented databases allow for better data compression ratios compared to row-oriented databases.

This is because compression algorithms can take advantage of the similar or repetitive values stored in each column.

For example, dictionary encoding can be applied to store unique values in a dictionary, replacing the original values with shorter references.

Common algorithms

Bit-Packing (small-range numeric values)

Huffman Coding (frequent categorical strings)

Byte Dictionary (repeated values, fewer than 256)


Bit Packing

Reduce storage by only using the minimum number of bits needed to represent values.

Suitable for columns with small ranges.

Example:

If temperatures are between 0 and 15, only 4 bits are needed instead of 32 (int)



Huffman Coding

Assigns shorter codes to frequently occurring values and longer codes to less frequent ones.

Suitable for columns with high-frequency categorical data, like product IDs.

Example: Product IDs that frequently appear in a sales table can be represented with shorter binary codes.

Shorter codes are assigned to more frequent IDs, Longer codes are assigned to less frequent IDs



Product ID	Frequency		Compressed
111111	100	→	0
222222	60		10
333333	30		110
444444	10		111


Byte Dictionary Encoding

Assigns unique byte values (0-255) to frequently occurring strings or values, storing each as a single byte.

Efficient for columns with many repeated values, such as product IDs or categories.

Example: Each product ID is assigned a unique byte, reducing storage size for frequent IDs.

Reverse Dictionary = each byte value in encoded_data maps back to the original product ID using reverse_dict



Illustrative Results

Average Temperature (Row Storage)	: 36.9554°C
Query Time (Row Storage)	: 17.57 milliseconds
Query Time (Columnar Storage)	: 5.13 milliseconds
Query time on compressed data	: 4.08 milliseconds

Illustrative Results

Total quantity sold: 412712, Total sales: \$2.26383e+07	
Uncompressed query time	: 7.79113 milliseconds
Huffman compressed query time	: 2.24537 milliseconds
Byte Dictionary compressed query time	: 0.707792 milliseconds

Vectorized Processing

standard for speeding up data processing on modern hardware
for building highly efficient analytical query engines

requires columnar representation of the data to write highly optimized query processing algorithms

significant departure from conventional tuple-based query processing model

requires a rewrite if query processing algorithms to work on columns

makes efficient utilization of CPU cache



instead of pushing one tuple at a time up the query plan tree, you basically push a block

difference between the vectorized model and the traditional model

block comprises of fixed set of tuples (records) represented by a set of vectors with one-to-one correspondence between vector and column/field in the schema

modern processors have extended instruction sets that extend the concept of vectorized execution to multiple column values in parallel in a single instruction – Single Instruction Multiple Data (SIMD)

SIMD exploits data level parallelism independent of concurrency.



Columnar storage

Star schemas work extremely well in columnar databases.

Better than 3NF designs.

To effectively use columnar storage in a star schema data warehouse, combine the structural advantages of the star schema with the performance benefits of columnar databases.

Optimize Queries:

Use column-specific selections, filters, and aggregations

Write Overhead:

Columnar storage is slower for row-wise inserts. Batch writes and use micro-batching in ETL pipelines

Storage Redundancy:

Star schemas duplicate dimension keys in fact tables, but columnar compression mitigates storage costs

Hybrid Approaches - Wide Tables for Reporting:

Create denormalized "One Big Table" (OBT) layers atop the star schema for self-service analytics, reducing join complexity

Use columnar storage for these tables to maintain query performance

But, because columnar databases have efficient compression and a lot of advantages for locating data in large tables, doesn't mean you throw data modeling best practices out the window.

You do not gain anything by implementing 'big flat tables' other than, maybe, shaving a few milliseconds off the query time.

Data modelling is mainly about providing organization and understanding to the data.

Parquet

Parquet file is an example of hybrid structure:

Row Groups: Horizontal split into logical blocks (e.g. 500 000 rows).

Column Chunks: Inside Row Groups data is split vertically into columns.

Pages: Smallest write unit (around 1 MB).

Here is the real encoding and compression, here the SQL engine is working on the data

File Footer: Contains schema information and location (addresses) of all column and row groups.


Typically contains also basic statistics, like min/max, null counts, for each page and row group.

Unit of parallelization:

MapReduce - File/Row Group

IO - Column chunk

Encoding/Compression - Page



Parquet

Encoding techniques:

Applied before Parquet compresses data (Snappy)

Dictionary Encoding:

Infile dictionary

Run-Length Encoding (RLE):

If data, after sorting, have a tendency to form long sequences

Bit Packing:

Knowing max/min values we can apply better representation

For instance to store values 0-3, we can use 2 bits and not full blown 32bit Integer.

Data	['Bruce', 'Bruce', 'Bruce', 'Bruce', 'Bruce', 'Bruce', 'Bruce', 'Clark', 'Clark', 'Clark', 'Clark', 'Clark', 'Clark']	
Dictionary Encoding	dictionary = ['Bruce', 'Clark']	replaces repeated values with shorter, unique keys
	indices = [0, 0, 0, 0, 0, 0, 0, 1, 1, 1, 1, 1]	
RLE Encoding	dictionary = ['Bruce', 'Clark']	the value 0 appears 7 times consecutive, and value 1 appears 5 times consecutive
	indices = [7, 0], [6, 1]	

Open-table formats

Parquet has one main fundamental flaw:

it is a single file format

Typically, we have table(s) comprised of many different Parquet files

How to manage such a structure

How to optimize their usage

In order to solve this problem **Open-Table Formats** were introduced

Provide programable metadata layer, which hides before the user the fact that a collection of Parquet (or other) files are just one consistent table

Provide SQL access, ACID, etc.

Tabular formats handle this issue by introducing a central log

Transaction Log / Metadata Layer

Open-table formats

Three main solutions are currently available on the market:

Delta Lake (Databricks)

Works based on a centralized transaction log (Transaction Log) located in the `_delta_log/` folder.

Every write operation creates a JSON file (e.g., `0001.json`), which states: "From now on, files A and B belong to the table, and file C is removed."

Advantages: *Exceptionally fast within the Apache Spark/Databricks ecosystem.*

Apache Iceberg (Netflix)

Iceberg takes a more hierarchical approach and relies on a so-called "metadata tree" – it defines a table through a tree of manifest files (text files describing other files), thereby not utilizing a folder structure.

Advantages: *It handles petabyte-scale tables because the engine does not need to scan the directory structure in the cloud (so-called object storage file listing).*

Apache Hudi (Uber)

Designed specifically for real-time data streaming and frequent updates (low-latency streaming).

Offers advanced mechanisms such as Merge-on-Read (updates are saved in small side-files and merged with the main Parquet file only when the table is read by the user).

Modern ROLAP

Modern storage architecture

Logical layer (Open-Tables - Delta/Iceberg):

Guarantees ACID, allows for Time Travel and makes sure that the BI system sees a consistent view of the data

File layer (Apache Parquet):

Responsible for ultraefficient arrangement of data into columns, their proper encoding and micro-indexing

Physical layer (Cloud Object Storage):

Cheap cloud storage, which physically stores data files

Or at least wants to – as underneath there is a pure key-value store

As such maintaining hierarchical structures of directories is not that easy



Partitioning

Partitioning or fragmentation

divides a table into smaller data sets to better support the management and processing of very large volumes of data

each one called a partition.

can be applied to tables as well as to indexes.

a partitioned index can be defined over an unpartitioned table

a partitioned table may have unpartitioned indexes defined over it.


There are two ways of partitioning a table: vertically and horizontally.

Vertical partitioning

splits the attributes of a table into groups that can be independently stored.

a table can be partitioned such that the most often used attributes are stored in one partition, while other less often used attributes are kept in another partition.

more records can be brought into main memory, reducing their processing time.




Horizontal partitioning

also known as database sharding

divides a table into smaller tables that have the same structure than the full table but fewer records.

if some queries require the most recent data while others access older data, a fact table can be horizontally partitioned according to some time frame

advantage is that refreshing the data warehouse is more efficient since only the last partition must be accessed.



Partitioning into smaller data sets

facilitates administrative tasks

increases query performance

especially when parallel processing is applied

enables access to a smaller subset of the data

if the user's selection does not refer to all partitions

is recommended when the contents of a table need to be distributed across different servers

advised in order to perform maintenance on parts of the data without invalidating the entire index

There are two classic techniques of partitioning related to query evaluation.

Partition pruning

is the typical way of improving query performance using partitioning, often producing performance enhancements of several orders of magnitude.


Partition-Wise Joins

The execution time of joins can also be enhanced by using partitioning.


break a large join into smaller joins that occur between each of the partitions

offers significant performance benefits both for serial and parallel execution.


This occurs when the two tables to be joined are partitioned on the join attributes or, in the case of foreign key joins, when the reference table is partitioned on its primary key.



With partition pruning




Range Partitioning



Hash Partitioning



List Partitioning



There are three most common partitioning strategies in database systems:

Range partitioning

which maps records to partitions based on ranges of values of the partitioning key.

The temporal dimension is a natural candidate for range partitioning, although other attributes can be used.

For example, if a table contains a date column defined as the partitioning key, the January 2012 partition will contain rows with key values from January 1, 2012, to January 31, 2012.

Hash partitioning

maps records to partitions based on a hashing algorithm applied to the partitioning key

The hashing algorithm distributes rows among partitions in a uniform fashion, yielding, ideally, partitions of the same size.

This is typically used when partitions are distributed in several devices and, in general, when data are not partitioned based on time since it is more likely to yield even record distribution across partitions.

List partitioning

enables to explicitly control how rows are mapped to partitions specifying a list of values for the partitioning key.





In this way, data can be organized in an ad hoc fashion.

Round-robin

Data is cyclically and equally distributed among partitions – regardless of the properties of the data, each split is sequentially assigned the next accessible data item.

Composite partitioning

Partitioning based on two or more partitioning techniques, e.g., range-range, range-hash.



Partitioning Technique	Description	Suitable Data	Query Performance	Data Distribution	Complexity
Horizontal Partitioning	Divides dataset based on rows/ records	Large datasets	Complex joins	Uneven distribution	Distributed transaction management
Vertical Partitioning	Divides dataset based on columns/ attributes	Wide tables	Improved retrieval	Efficient storage	Increased query complexity
Key-based Partitioning	Divides dataset based on specific key	Key-value datasets	Efficient key lookups	Even distribution by key	Limited query flexibility
Range Partitioning	Divides dataset based on specific range	Ordered datasets	Efficient range queries	Even distribution by range	Joins and range queries
Hash-based Partitioning	Divides dataset based on hash function	Unordered datasets	Even distribution	Random distribution	Inefficient key-based queries
Round-robin Partitioning	Divides dataset in a cyclic manner	Equal-sized datasets	Basic load balancing	Even distribution	Limited query optimization

MOLAP

Partitioned Cubes

To overcome the space limitation of MOLAP, the cube can be partitioned

One logical cube of data can be spread across multiple physical cubes on separate (or same) servers

The divide&conquer cube partitioning approach helps alleviate the scalability limitations of MOLAP implementation


Ideal cube partitioning is completely invisible to end users

Performance degradation does occur in case of a join across partitioned cubes

MOLAP

Advanced MOLAP implementations

Allow to distribute queries across partitioned cubes (potentially on distinct server platforms) for parallel retrieval



Jensen C.S., Pedersen T.B., Thomsen C.,
Multidimensional Databases and Data Warehousing,
Morgan & Claypool Publishers series SYNTHESIS LECTURES ON DATA MANAGEMENT,
2010

Inmon W.,
Building the Data Warehouse,
John Wiley & Sons, New York 2002

Sherman R.,
Business Intelligence Guidebook
Elsevier, 2015

Claudia Imhoff, Nicholas Gallempo, Jonathan G. Geiger,
Mastering Data Warehouse Design - Relational and Dimensional Techniques,
Wiley Publishing, Inc., 2003

Efraim Turban, Ramesh Sharda, Dursun Delen, David King, Janine E. Aronson
Business Intelligence: A Managerial Approach
Prentice Hall, 2010

Sanjiv Purba (ed.)
Handbook of Data Management (1999 Edition)
Auerbach, 1999

Joy Arulraj
Lecture for Database System Implementation, 2024
Georgia Tech, College of Computing
<https://faculty.cc.gatech.edu/~jarulraj/courses/4420-f24/slides/22-columnar-storage-and-compression.pdf>

Bibliography

SOURCES



View maintenance

Updating cuboids

In data warehouses, materialized views that include aggregate functions are called **summary tables**

The problem:

as data at the sources are added or updated, the summary tables that depend on these data must be also updated

summary tables can be maintained with minimum access to the base data while keeping maximum data availability

Two options:

to recompute the summary tables from scratch

to apply incremental view maintenance techniques to avoid such recomputation

View maintenance

under certain conditions, it is possible to compute changes to the view caused by changes in the underlying relations without recomputing the entire view from scratch.

incremental view maintenance

Sales(ProductKey, CustomerKey, Quantity)

TopProducts = $\pi_{\text{ProductKey}}(\sigma_{\text{Quantity} > 150}(\text{Sales}))$

 (p2, c3, 110)

 (p2, c3, 160)

 (p2, c3, 160) 

requires to scan the
relation Sales

view maintenance

under certain conditions, it is possible to compute changes to the view caused by changes in the underlying relations without recomputing the entire view from scratch.

incremental view maintenance

Sales(ProductKey, CustomerKey, Quantity)

FoodCustomers = $\pi_{\text{CustomerKey}}(\sigma_{\text{CategoryName}='Food'}(\text{Product}) * \text{Sales})$

 (p3, c4, 170) 

we check in the base relations whether or not p3 is in the food category

An aggregate function is **self-maintainable**

if the new value of the function can be computed solely from the old values and from changes to the base data

Aggregate functions must be distributive in order to be self-maintainable.

The five classic aggregate functions in SQL are self-maintainable with respect to insertions, but not to deletions.

SUM, AVG, COUNT, MAX, MIN

In fact, MAX and MIN are not, and cannot be made, self-maintainable with respect to deletions

The summary-delta algorithm

a representative summary table maintenance algorithm
has two main phases called propagate and refresh:

*propagate phase can be performed in parallel with data warehouse operations
only the refresh phase requires taking the warehouse off-line*

basic idea is to

*create in the propagate phase a so-called **summary-delta table** that stores the net changes to the summary table due to changes in the source data
during the refresh phase, these changes are applied to the summary table*

Idea:

processes two sets of deferred changes:

pos-ins: Table storing deferred insertions; Projects inserted tuples with positive values (e.g. +1 for count, +qty for quantity)

pos-del: Table storing deferred deletions; Projects deleted tuples with negative values (e.g., -1 for count, -qty for quantity)

*later unions both results and aggregates by the same group-by attributes as the summary table
produces net changes for each aggregate function value*

DailySalesSum

```
CREATE VIEW DailySalesSum AS (  
    SELECT    ProductKey, TimeKey, SUM(Quantity) AS SumQuantity,  
              COUNT(*) AS Count  
    FROM      Sales  
    GROUP BY ProductKey, TimeKey )
```

In the propagate phase, we define two tables

$\Delta^+(\text{Sales})$ and $\Delta^-(\text{Sales})$

store the insertions and deletions to the fact table,

and a view where the net changes to the summary tables are stored

```
CREATE VIEW SD_DailySalesSum(ProductKey, TimeKey,  
    SD_SumQuantity, SD_Count) AS  
    WITH Temp AS (  
        ( SELECT ProductKey, TimeKey,  
              Quantity AS _Quantity, 1 AS _Count  
        FROM     $\Delta^+(\text{Sales})$  )  
        UNION ALL  
        ( SELECT ProductKey, TimeKey,  
              -1 * Quantity AS _Quantity, -1 AS _Count  
        FROM     $\Delta^-(\text{Sales})$  ) )  
    SELECT ProductKey, TimeKey, SUM(_Quantity), SUM(_Count)  
    FROM Temp  
    GROUP BY ProductKey, TimeKey
```

DailySalesSum

```
CREATE VIEW DailySalesSum AS (  
    SELECT    ProductKey, TimeKey, SUM(Quantity) AS SumQuantity,  
              COUNT(*) AS Count  
    FROM      Sales  
    GROUP BY ProductKey, TimeKey )
```

During the refresh phase, we apply to the summary table the net changes stored in the summary-delta table

```
For each tuple T in SD_DailySalesSum DO  
    IF NOT EXISTS (  
        SELECT *  
        FROM    DailySalesSum D  
        WHERE   T.ProductKey = D.ProductKey AND  
                T.TimeKey = D.TimeKey)  
        INSERT T INTO DailySalesSum  
    ELSE  
        IF EXISTS (  
            SELECT *  
            FROM    DailySalesSum D  
            WHERE   T.ProductKey = D.ProductKey AND  
                    T.TimeKey = D.TimeKey AND  
                    T.SD_Count + D.Count = 0)  
            DELETE T FROM DailySalesSum  
        ELSE  
            UPDATE DailySalesSum  
            SET SumQuantity = SumQuantity + T.SD_SumQuantity,  
                Count = Count + T.SD_Count  
            WHERE ProductKey = T.ProductKey AND  
                  TimeKey = T.TimeKey
```

DailySalesSum

Product Key	Time Key	Sum Quantity	Count
p2	t2	100	10
p3	t3	100	20
p4	t4	100	15
p5	t5	100	2
p6	t6	100	1

$\Delta^+(\text{Sales})$

Product Key	Customer Key	Time Key	Quantity
p1	c1	t1	150
p2	c1	t2	200
p2	c2	t2	100
p4	c2	t4	100
p6	c5	t6	200

$\Delta^-(\text{Sales})$

Product Key	Customer Key	Time Key	Quantity
p2	c1	t2	10
p5	c2	t5	10
p6	c5	t6	100

SD_DailySalesSum

Product Key	Time Key	SD_Sum Quantity	SD_Count
p1	t1	150	1
p2	t2	290	1
p4	t4	100	1
p5	t5	-10	-1
p6	t6	100	0

DailySalesSum after update

Product Key	Time Key	Sum Quantity	Count
p1	t1	150	1
p2	t2	390	11
p3	t3	100	20
p4	t4	200	16
p5	t5	90	1
p6	t6	200	1

Product Key	Time Key	Sum Quantity	Count
p2	t2	100	10
p3	t3	100	20
p4	t4	100	15
p5	t5	100	2
p6	t6	100	1



Product Key	Time Key	Sum Quantity	Count
p1	t1	150	1
p2	t2	390	11
p3	t3	100	20
p4	t4	200	16
p5	t5	90	1
p6	t6	200	1

Product Key	Time Key	SD_Sum Quantity	SD_Count
p1	t1	150	1
p2	t2	290	1
p4	t4	100	1
p5	t5	-10	-1
p6	t6	100	0

Product Key	Customer Key	Time Key	Quantity
p1	c1	t1	150
p2	c1	t2	200
p2	c2	t2	100
p4	c2	t4	100
p6	c5	t6	200



Product Key	Customer Key	Time Key	Quantity
p2	c1	t2	10
p5	c2	t5	10
p6	c5	t6	100



DailySalesMax

Product Key	Time Key	Max Quantity	Count
p2	t2	150	4
p3	t3	100	5
p4	t4	50	6
p5	t5	150	5
p6	t6	100	3
p7	t7	100	2

SD_DailySalesMax

Product Key	Time Key	SD_Max Quantity	SD_Count
p1	t1	100	2
p2	t2	100	2
p3	t3	100	2
p4	t4	100	2
p5	t5	100	-2
p6	t6	100	-2
p7	t7	100	-2

Updated view DailySalesMax


Product Key	Time Key	Max Quantity	Count
p1	t1	100	2
p2	t2	150	4
p3	t3	100	7
p4	t4	100	2
p5	t5	150	5
p6	t6	100	1
p7	t7	?	?

Product Key	Time Key	Max Quantity	Count
p2	t2	150	4
p3	t3	100	5
p4	t4	50	6
p5	t5	150	5
p6	t6	100	3
p7	t7	100	2

Product Key	Time Key	Max Quantity	Count
p1	t1	100	2
p2	t2	150	4
p3	t3	100	7
p4	t4	100	2
p5	t5	150	5
p6	t6	100	1
p7	t7	?	?



Product Key	Time Key	SD_Max Quantity	SD_Count
p1	t1	100	2
p2	t2	100	2
p3	t3	100	2
p4	t4	100	2
p5	t5	100	-2
p6	t6	100	-2
p7	t7	100	-2




Data Computation

Data Cube Computation Methods

Data cube computation is an essential task in data warehouse implementation.

The precomputation of all or part of a data cube can greatly reduce the response time and enhance the performance of online analytical processing.

However, such computation is challenging because it may require substantial computational time and storage space.



Partial order

$Q1=(\text{part}, \text{supplier}, *)$ can be derived based on $Q2=(\text{part}, \text{supplier}, \text{customer})$

$$Q1 < Q2$$

Also for hierarchies

$$(\text{part}, \text{country}, *) < (\text{part}, \text{city}, *)$$

With a constrained space of 2^S (the power space of a set of views), what is the optimal choice of views to materialize?

Cost model

Estimating the benefits of materializing a view – related to the evaluation strategy

It is often impossible to know the sizes of views without actually calculating them

Data:

for each view, its size, the number of k

Search:

k views with the highest profit

It is a NP-Complete problem

The problem of optimally selecting cuboids to materialize in a data warehouse can be reduced to the **set cover problem**

Algorithms like Bottom-Up Computation and Star-Cubing achieve $O(n \log n)$ for specific cases

Several algorithmic strategies have been developed to address the view materialization optimization problem:

Greedy Algorithms:

These approaches iteratively select views that provide the greatest benefit-to-cost ratio, considering factors such as query frequency, view size, and maintenance requirements.

The greedy method incorporates maintenance costs and storage constraints while utilizing "And-Or" view graphs to represent all possible ways to generate warehouse views.

Heuristic Algorithms:

These methods utilize Multiple View Processing Plans (MVPP) to obtain optimal materialized view selections that achieve the best combination of performance and low maintenance cost.

Heuristic approaches can provide feasible solutions based on individual optimal query plans when exact optimization becomes computationally prohibitive.

Genetic Algorithms:

Advanced evolutionary approaches have been developed that use binary encoding where 0 indicates a view is not materialized and 1 indicates materialization.

These algorithms have shown superior performance compared to conventional heuristic methods in terms of both speed and accuracy.

Selection Criteria and Best Practices for optimal view materialization requires careful consideration of several key factors:

Query Patterns and Frequency:

Views that support frequently executed queries should receive higher priority for materialization, as the performance benefits will be realized more often.

Analysis of historical query workloads helps identify the most valuable materialization candidates.

Data Filtering and Aggregation:

The most effective materialized views typically perform one or both of two key functions: filtering data (either by rows or columns) and performing resource-intensive operations such as complex aggregations.

Views that filter to recent data or unusual patterns often provide significant performance benefits.

Storage Effectiveness:

The concept of storage effectiveness (η) measures the benefit-to-storage ratio for each potential materialized view.

Views with high storage effectiveness yield the greatest performance improvement per unit of storage space consumed.


Maintenance Windows:

The available time for view maintenance operations constrains which views can be practically materialized.

Views requiring frequent updates may not be suitable for environments with limited maintenance windows.


Greedy approach

Query	Size
a	100
b	50
c	75
d	20
e	30
f	40
g	1
h	10



Greedy approach


Query	Size
a	100
b	50
c	75
d	20
e	30
f	40
g	1
h	10



Query	Size	Uses?	Cost	Benefit
a	100	a	100	-
b	50	b	50	50
c	75	a	100	-
d	20	b	50	50
e	30	b	50	50
f	40	a	100	-
g	1	b	50	50
h	10	b	50	50

Greedy approach


Query	Size
a	100
b	50
c	75
d	20
e	30
f	40
g	1
h	10



Q	b	c	d	e	f	g	h
a							
b	50						
c		25					
d	50		80				
e	50	25		70			
f		25			60		
g	50	25	80	70		99	
h	50	25		70	60		90
Total	250	125	160	210	120	99	90

Greedy approach


Query	Size
a	100
b	50
c	75
d	20
e	30
f	40
g	1
h	10



Q	c	d	e	f	g	h
a						
b						
c	25					
d		30				
e			20			
f	25			60		
g		30	20		49	
h			20	60		40
Total	50	60	60	70	49	40

Greedy approach

Query	Size
a	100
b	50
c	75
d	20
e	30
f	40
g	1
h	10



Q	c	d	e	g	h
a					
b					
c	25				
d		30			
e			20		
f					
g		30	20	49	
h			20		30
Total	25	60	60	49	30

There are general optimization techniques for efficient computation of data cubes.

Optimization Technique 1: Sorting, hashing, and grouping.

In cube computation, aggregation is performed on the tuples (or cells) that share the same set of dimension values.

it is important to explore sorting, hashing, and grouping operations to access and group such data together to facilitate computation of such aggregates.

For example,

to compute total sales by branch, day, and item, it can be more efficient to sort tuples or cells by branch, and then by day, and then group them according to the item name.

This technique can also be further extended to perform

shared-sorts (i.e., sharing sorting costs across multiple cuboids when sort-based methods are used),

shared-partitions (i.e., sharing the partitioning cost across multiple cuboids when hash-based algorithms are used).

There are general optimization techniques for efficient computation of data cubes.

Optimization Technique 2: Simultaneous aggregation and caching intermediate results

It is efficient to compute higher level aggregates from previously computed lower-level aggregates

rather than from the base fact table

Simultaneous aggregation from cached intermediate computation results may lead to the reduction of expensive disk I/O operations

For example,

to compute sales by branch, we can use the intermediate results derived from the computation of a lower-level cuboid, such as sales by branch and day.

This technique can be further extended to perform amortized scans

i.e., computing as many cuboids as possible at the same time to amortize disk reads

There are general optimization techniques for efficient computation of data cubes.

Optimization Technique 3: Aggregation from the smallest child, when there exist multiple child cuboids.

When there exist multiple child cuboids

it is usually more efficient to compute the desired parent (i.e., more generalized) cuboid from the smallest, previously computed child cuboid.

For example to compute a sales cuboid – C_{branch}

when there exist two previously computed cuboids

$C_{\{branch, year\}}$ and $C_{\{branch, item\}}$

it is obviously more efficient to compute C_{branch} from the former than from the latter

if there are many more distinct items than distinct years

There are general optimization techniques for efficient computation of data cubes.

Optimization Technique 4: The Apriori pruning method can be explored to compute iceberg cubes efficiently.

The Apriori property, in the context of data cubes, states as follows: If a given cell does not satisfy minimum support, then no descendant of the cell (i.e., more specialized cell) will satisfy minimum support either

can be used to substantially reduce the computation of iceberg cubes.

A common iceberg condition is that the cells must satisfy a minimum support threshold, such as a minimum count or sum

the Apriori property can be used to prune away the exploration of the descendants of the cell

if a condition is violated for some cell c , then every descendant of c will also violate that condition

For example, if the count of a cell, c , in a cuboid is less than a minimum support threshold, v , then the count of any of c 's descendant cells in the lower-level cuboids can never be greater than or equal to v , and thus can be pruned.

Computing cubes:

3 methodologies

Bottom-Up:

Multi-Way array aggregation (Zhao, Deshpande & Naughton, SIGMOD'97)

Top-down:

Bottom-up cube computation – BUC (Beyer & Ramakrishnan, SIGMOD'99)

H-cubing technique (Han, Pei, Dong & Wang: SIGMOD'01)

Integrating Top-Down and Bottom-Up:

Star-cubing algorithm (Xin, Han, Li & Wah: VLDB'03)

High-dimensional OLAP:

A Minimal Cubing Approach (Li, et al. VLDB'04)